

Arc-Pico-VL-v0.1

15M-class cheap multimodal pretraining testbed

15.27M
Total

13.70M
Text


1.57M
Vision

512
Context

Arc-Pico is intentionally tiny. It is meant to prove tokenizer, data, mixed text/image training, checkpointing, and image-use validation before Arc-124M spends serious GPU money.

Model Flow

Two input paths, one shared decoder stream



The image becomes 16 soft tokens. Those tokens have the same 256-wide shape as text embeddings, so the decoder can attend over one mixed sequence.

Configuration

Implemented dimensions

Field	Value
vocab_size	16384
d_model	256
n_layers	10
n_heads	8
head_dim	32
ffn_hidden	896
image_size	224
vision_width	192
vision_layers	2
vision_resampler_tokens	16

500M mixed tokens/equivalent
85% text only
Training target
15% image-caption
single GPU
under about 5 EUR target